

Course Title:	Advanced Computer Architecture
Course Number:	COE 818
Semester/Year (e.g.F2016)	W2023

Instructor:	Dr. Arghavan Asad
--------------------	-------------------

<i>Assignment/Lab Number:</i>	4
<i>Assignment/Lab Title:</i>	Using NVIDIA GPU

<i>Submission Date:</i>	March 13, 2023
<i>Due Date:</i>	March 13, 2023

Student LAST Name	Student FIRST Name	Student Number	Section	Signature*
Youssef	Helana	500766171	1	H.Y.

*By signing above you attest that you have contributed to this written lab report and confirm that all work you have contributed to this lab report is your own work. Any suspicion of copying or plagiarism in this work will result in an investigation of Academic Misconduct and may result in a "0" on the work, an "F" in the course, or possibly more severe penalties, as well as a Disciplinary Notice on your academic record under the Student Code of Academic Conduct, which can be found online at: <https://www.torontomu.ca/content/dam/senate/policies/pol60.pdf>

Lab 4: Using NVIDIA GPU

1. Objective:

We create a parallel matrix multiplication application in this lab. The matrix multiplication application is run on CUDA, and we discover that adding more threads improves speed.

2. Introduction:

NVIDIA created CUDA as a general-purpose parallel computing framework for use with its own GPUs. By utilising GPU capability, CUDA enables programmers to accelerate computationally demanding programmes. Several GPU APIs, like OpenCL, have been proposed. There are moreover other competitive GPUs from firms like AMD. Several application domains, including deep learning, are dominated by the CUDA and NVIDIA GPU combination [1].

We create a parallel matrix multiplication application in this lab. The matrix multiplication application is run on CUDA, and we discover that adding more threads improves speed. We assess the application's performance when the number of threads per block is changed from 2 to 32 in multiples of 2. We use square matrices of sizes 512 x 512 and 1024 x 1024 to compare the performance of the matrix multiplication application in each case.

3. Background

Using Device Query to Determine the Optimal Threads Per Block

It is essential to examine the characteristics of the CUDA device in the system that is used to run the application to choose the best number of threads per block for our matrix multiplication application. We make use of the Device Query example found in the folder /cuda-7.0/samples/1 Utilities/deviceQuery.

The CUDA devices that are present in the system are listed by their attributes in this example. The device characteristics of the system that we utilize to run the matrix multiplication application are shown in Figure 1.

```

cuda_shell
***** CUDA Work Environment *****

Welcome to the CUDA environment. The environment has
been setup for you to compile CUDA supported GPU programs.

Changing to existing cuda folder: /home/student1/h2yousse/cuda-7.0 ...Done
[h2yousse@bloor:~/cuda-7.0/samples]$ ~/cuda-7.0/samples/0_Simple/matrixMul
bash: /home/student1/h2yousse/cuda-7.0/samples/0_Simple/matrixMul: Is a director
y
[h2yousse@bloor:~/cuda-7.0/samples]$ cd 0_Simple
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple]$ cd matrixMul
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arc
h=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=comput
e_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,cod
e=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=comput
e_52 -o matrixMul.o -c matrixMul.cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,cod
e=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35
-gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode
arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMu
l matrixMul.o
mkdir -p ../../bin/x86_64/linux/release
cp matrixMul ../../bin/x86_64/linux/release
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ ^C
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/sample
s/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 512
==19238== NVPROF is profiling process 19238, command: /home/student1/h2yousse/cu
da-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 2
Grid Size X: 256 , Grid Size Y: 256
Elapsed Time for CPU Multiplication: 1.490000 sec
SUCCESS!
==19238== Profiling application: /home/student1/h2yousse/cuda-7.0/samples/0_Simp
le/matrixMul/matrixMul
==19238== Profiling result:
Time(%) Time Calls Avg Min Max Name
99.66% 595.90ms 1 595.90ms 595.90ms 595.90ms matrixMul(int*, int*,
int*, int)
0.23% 1.3484ms 2 674.20us 662.67us 685.74us [CUDA memcpy HtoD]
0.11% 671.27us 1 671.27us 671.27us 671.27us [CUDA memcpy DtoH]

==19238== API calls:
Time(%) Time Calls Avg Min Max Name
99.25% 528.36ms 3 176.12ms 255.63us 527.84ms cudaMalloc
0.55% 2.9424ms 3 980.80us 452.77us 1.4293ms cudaMemcpy
0.10% 524.28us 83 6.3160us 1.0100us 199.70us cuDeviceGetAttribute
0.06% 294.08us 3 98.026us 77.265us 136.26us cudaFree
0.01% 64.283us 1 64.283us 64.283us 64.283us cuDeviceTotalMem
0.01% 54.692us 1 54.692us 54.692us 54.692us cudaLaunch
0.01% 53.707us 1 53.707us 53.707us 53.707us cuDeviceGetName
0.01% 27.220us 4 6.8050us 1.5790us 14.233us cudaSetupArgument
0.00% 5.3010us 2 2.6500us 1.1950us 4.1060us cuDeviceGet
0.00% 4.3200us 2 2.1600us 1.5520us 2.7680us cuDeviceGetCount
0.00% 4.2440us 1 4.2440us 4.2440us 4.2440us cudaConfigureCall
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$

```

Figure 1: Properties of the CUDA Device in the System Used to Run the Matrix Multiplication Application

You can see that the matrix multiplication application is running on a CUDA device that uses 48 cores per multiprocessor in the device's characteristics. Also take note of the warp size, which is 32. We must thus make sure to utilise a number of threads per block that is a multiple of the warp size (32) and the number of cores per multiprocessor (48 cores per multiprocessor) when determining the ideal number of threads per block to run the matrix multiplication application [2].

As the number of threads per block is increased in the range of 2 to 16 threads per block, we anticipate the scalability of the matrix multiplication application to grow (in multiples of 2). We anticipate that the scalability of the matrix multiplication application would decline when the number of threads per block is increased from 16 to 32 threads per block since 32 threads per block is not a multiple of 48 cores per multiprocessor.

4. Methodology

We create a straightforward CUDA application for matrix multiplication. The matrix multiplication programme is executed one step at a time. The matrix multiplication application is then run in parallel with thread counts per block ranging from 2 to 32 in multiples of 2. We keep track of how long it takes to finish each matrix multiplication. Also, we keep track of how long it takes to perform matrix multiplications for each of the scenarios previously discussed using square matrices of sizes 512 x 512 and 1024 x 1024.

The appendix contains the code for the matrix multiplication application. In the results part of the appendix, Table I summarizes the results that were shown after the application was run.

5. Results

A) Total Time Needed to Finish Matrix Multiplication

Figure 2 and Table I both show the amount of time needed to complete each of the matrix multiplications. The results show that a considerable time improvement was achieved by raising the number of threads per block from 2 threads per block to 16 threads per block in multiples of 2.

Moreover, going from 16 to 32 threads per block did not speed up the process. This is because 32 threads per block is not a multiple of the 48 cores that our multiprocessor has. Hence, 16 threads per block is the ideal number of threads for our matrix multiplication application.

Table I: Elapsed Time to Complete Matrix Multiplication

Matrix Size	Time (ms)					
	CPU	GPU				
		Threads per Block = 2	Threads per Block = 4	Threads per Block = 8	Threads per Block = 16	Threads per Block = 32
512 x 512	1380	595.92	189.56	50.197	29.557	36.658
1024 x 1024	10240	5847.46	1548.97	406.21	247.64	288.56

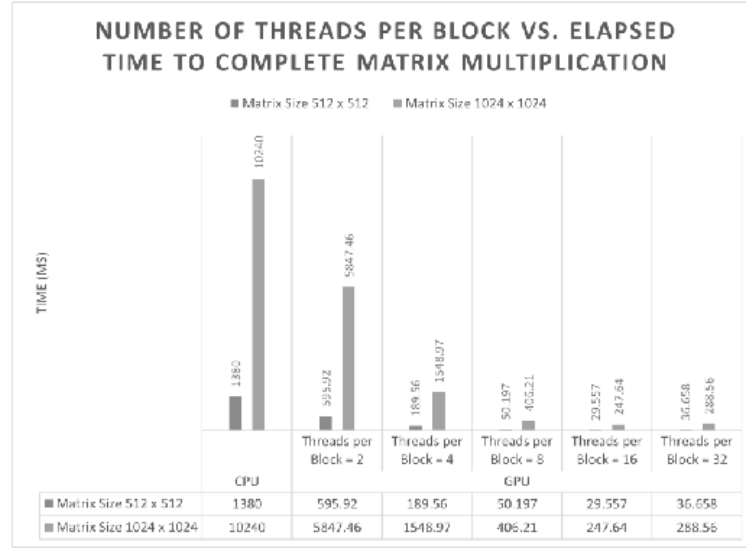


Figure 2: Number of Threads per Block vs. Elapsed Time to Complete Matrix Multiplication

B) Scalability

The scalability is determined by comparing the amount of time needed to multiply a matrix sequentially on a CPU to the amount of time needed to multiply the same matrix in parallel on a GPU utilising x threads per block, as shown in (1):

$$Scalability = \frac{T_{seq}}{T_{parallel}(x \text{ Threads per Block})} \quad (1)$$

For the situations of square matrices of size 512 x 512 and square matrices of size 1024 x 1024, Table II shows the scalability of the matrix multiplication application when the number of threads per block is increased from 2 to 32 in multiples of 2. The statistics presented in Table II are illustrated in Figure 3.

The results in Table II and Figure 3 make it clear that the matrix multiplication application is more scalable as the number of threads per block rises, from 2 to 16 threads per block, in multiples of 2. It should be noted that scaling the matrix multiplication application increases the number of threads per block from 16 to 32 but limits its scalability. This is because 32 threads per block is not a multiple of the 48 cores that our multiprocessor has. Hence, 16 threads per block is the ideal number of threads for our matrix multiplication application.

Table II: Scalability of Matrix Multiplication Application

Matrix Size	Scalability					
	CPU	GPU				
		Threads per Block = 2	Threads per Block = 4	Threads per Block = 8	Threads per Block = 16	Threads per Block = 32
512 x 512	1	2.315747	7.280017	27.49168	46.68945	37.64526
1024 x 1024	1	1.751188	6.610845	25.20864	41.35035	35.48655




Figure 3: Number of Threads per Block vs. Scalability of Matrix Multiplication Application

C) Time Overhead

The time overhead is calculated using (2):

$$T_{overhead} = T_{parallel}(Threads\ per\ Block) - \frac{T_{seq}}{Total\ Number\ of\ Threads} \quad (2)$$

As the number of threads per block increases from 2 to 16 threads per block, in multiples of 2, the matrix multiplication application becomes more scalable, as shown by the findings in Table II and Figure 3. It should be noted that scaling the matrix multiplication application results in a 16 to 32 thread increase per block but limits its potential to scale. This is because our multiprocessor only has 48 cores, although there are 32 threads per block. This means that the ideal number of threads for our matrix multiplication application is 16 threads per block.

Table III: Time Overhead of Matrix Multiplication Application

Matrix Size	Time Overhead				
	Threads per Block = 2	Threads per Block = 4	Threads per Block = 8	Threads per Block = 16	Threads per Block = 32
512 x 512	595.909	189.539	50.1549	29.4728	36.4895
1024 x 1024	5847.44	1548.93	406.132	247.484	288.248




Figure 4: Number of Threads per Block vs. Time Overhead of Matrix Multiplication Application

6. Conclusion

The statistics in Table II and Figure 3 clearly show that as the number of threads per block rises from 2 to 16, the scalability of the matrix multiplication application also increases. It should be noted that scaling the matrix multiplication application increases the number of threads per block from 16 to 32 but limits its scalability. This is because 32 threads per block is not a multiple of the 48 cores that our multiprocessor has. Hence, 16 threads per block is the ideal number of threads for our matrix multiplication application.

The time overhead of the matrix multiplication application is greatest when 2 threads per block are used, as can be seen from the data shown in Table III and shown in Figure 4. When the number of threads per block is changed, in multiples of 2, from 2 to 16 threads per block, the matrix multiplication application's time overhead decreases. When 16 threads per block are used, the matrix multiplication application's time overhead is at its lowest. It should be noted that the time overhead grows as the number of threads per block rises from 16 to 32. This is because 32 threads per block is not a multiple of the 48 cores that our multiprocessor has.

7. References

- [1] M. Heller, "What is CUDA? Parallel Programming for GPUs," InfoWorld, 2018.
<https://www.infoworld.com/article/3299703/what-is-cuda-parallel-programming-for-gpus.html#:~:text=CUDA is a parallel computing, parallelizable part of the computation.>
- [2] C. Frederick, "The optimal number of threads per block in CUDA programming," Research Gate, 2017.
https://www.researchgate.net/post/The_optimal_number_of_threads_per_block_in_CUDA_programming.

8. Appendix

A) Console Results

```
cuda_shell
***** CUDA Work Environment *****

Welcome to the CUDA environment. The environment has
been setup for you to compile CUDA supported GPU programs.

Changing to existing cuda folder: /home/student1/h2yousse/cuda-7.0 ...Done
[h2yousse@bloor:~/cuda-7.0/samples]$ ~/cuda-7.0/samples/0_Simple/matrixMul
bash: ~/cuda-7.0/samples/0_Simple/matrixMul: Is a director
y
[h2yousse@bloor:~/cuda-7.0/samples]$ cd 0_Simple
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple]$ cd matrixMul
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul.o -c matrixMul.cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul matrixMul.o
mkdir -p ../../bin/x86_64/linux/release
cp matrixMul ../../bin/x86_64/linux/release
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ ^C
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 512
==19238== NVPROF is profiling process 19238, command: /home/student1/h2yousse/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 2
Grid Size X: 256 , Grid Size Y: 256
Elapsed Time for CPU Multiplication: 1.490000 sec
SUCCESS!
==19238== Profiling application: /home/student1/h2yousse/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
==19238== Profiling result:
Time(%) Time Calls Avg Min Max Name
99.66% 595.90ms 1 595.90ms 595.90ms 595.90ms matrixMul(int*, int*, int*, int)
0.23% 1.3484ms 2 674.20us 662.67us 685.74us [CUDA memcpy HtoD]
0.11% 671.27us 1 671.27us 671.27us 671.27us [CUDA memcpy DtoH]

==19238== API calls:
Time(%) Time Calls Avg Min Max Name
99.25% 528.36ms 3 176.12ms 255.63us 527.84ms cudaMalloc
0.55% 2.9424ms 3 980.80us 452.77us 1.4293ms cudaMemcpy
0.10% 524.28us 83 6.3160us 1.0100us 199.70us cuDeviceGetAttribute
0.06% 294.08us 3 98.026us 77.265us 136.26us cudaFree
0.01% 64.283us 1 64.283us 64.283us 64.283us cuDeviceTotalMem
0.01% 54.692us 1 54.692us 54.692us 54.692us cudaLaunch
0.01% 53.707us 1 53.707us 53.707us 53.707us cuDeviceGetName
0.01% 27.220us 4 6.8050us 1.5790us 14.233us cudaSetupArgument
0.00% 5.3010us 2 2.6500us 1.1950us 4.1060us cuDeviceGet
0.00% 4.3200us 2 2.1600us 1.5520us 2.7680us cuDeviceGetCount
0.00% 4.2440us 1 4.2440us 4.2440us 4.2440us cudaConfigureCall
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$
```

Figure 5: Console Results for Matrix Multiplication Application Using Matrices of Size 512 x 512 and 2 Threads per Block

```

[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -g
cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -g
mkdir -p ../bin/x86_64/linux/release
cp matrixMul ../bin/x86_64/linux/release
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 512
==19611== NVPF is profiling process 19611, command: /home/student1/h2yousse/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 4
Grid Size X: 128 , Grid Size Y: 128
Elapsed Time for CPU Multiplication: 1.440000 sec
SUCCESS!
==19611== Profiling application: /home/student1/h2yousse/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
==19611== Profiling result:
Time(%) Time Calls Avg Min Max Name
98.97% 189.09ms 1 189.09ms 189.09ms 189.09ms matrixMul(int*, int*, int*, int)
0.70% 1.3312ms 2 665.61us 662.67us 668.55us [CUDA memcpy HtoD]
0.34% 644.14us 1 644.14us 644.14us 644.14us [CUDA memcpy DtoH]

==19611== API calls:
Time(%) Time Calls Avg Min Max Name
96.43% 100.02ms 3 33.340ms 156.29us 99.700ms cudaMalloc
2.52% 2.6102ms 3 870.08us 291.01us 1.3981ms cudaMemcpy
0.60% 619.22us 83 7.4600us 1.1830us 249.93us cuDeviceGetAttribute
0.27% 275.61us 3 91.870us 78.584us 114.81us cudaFree
0.07% 77.208us 1 77.208us 77.208us 77.208us cuDeviceTotalMem
0.06% 63.638us 1 63.638us 63.638us 63.638us cuDeviceGetName
0.03% 34.177us 1 34.177us 34.177us 34.177us cudaLaunch
0.01% 10.736us 4 2.6840us 948ns 7.4070us cudaSetupArgument
0.01% 6.1090us 2 3.0540us 1.3340us 4.7750us cuDeviceGet
0.00% 5.0900us 2 2.5450us 1.8370us 3.2530us cuDeviceGetCount
0.00% 2.7360us 1 2.7360us 2.7360us 2.7360us cudaConfigureCall
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$

```

Figure 6: Console Results for Matrix Multiplication Application Using Matrices of Size 512 x 512 and 4 Threads per Block

```

[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul.o -c matrixMul.cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul matrixMul.o
mkdir -p ../bin/x86_64/linux/release
cp matrixMul ../bin/x86_64/linux/release
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 512
==19908== NVPF is profiling process 19908, command: /home/student1/h2yousse/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 8
Grid Size X: 64 , Grid Size Y: 64
Elapsed Time for CPU Multiplication: 1.470000 sec
SUCCESS!
==19908== Profiling application: /home/student1/h2yousse/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
==19908== Profiling result:
Time(%) Time Calls Avg Min Max Name
96.17% 50.132ms 1 50.132ms 50.132ms 50.132ms matrixMul(int*, int*, int*, int)
2.56% 1.3328ms 2 666.41us 662.70us 670.12us [CUDA memcpy HtoD]
1.27% 661.45us 1 661.45us 661.45us 661.45us [CUDA memcpy DtoH]

==19908== API calls:
Time(%) Time Calls Avg Min Max Name
95.87% 82.949ms 3 27.650ms 153.09us 82.633ms cudaMalloc
3.07% 2.6577ms 3 885.91us 288.17us 1.4462ms cudaMemcpy
0.52% 452.58us 83 5.4520us 995ns 164.25us cuDeviceGetAttribute
0.34% 294.63us 3 98.211us 76.981us 137.36us cudaFree
0.07% 57.044us 1 57.044us 57.044us 57.044us cuDeviceTotalMem
0.05% 45.961us 1 45.961us 45.961us 45.961us cuDeviceGetName
0.05% 44.318us 1 44.318us 44.318us 44.318us cudaLaunch
0.02% 13.103us 4 3.2750us 1.1300us 9.0310us cudaSetupArgument
0.01% 4.8560us 2 2.4280us 1.1890us 3.6670us cuDeviceGet
0.00% 3.9480us 2 1.9740us 1.3930us 2.5550us cuDeviceGetCount
0.00% 3.2750us 1 3.2750us 3.2750us 3.2750us cudaConfigureCall
[h2yousse@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$

```

Figure 7: Console Results for Matrix Multiplication Application Using Matrices of Size 512 x 512 and 8 Threads per Block

```
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul.o -c matrixMul.cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul matrixMul.o
mkdir -p ../bin/x86_64/linux/release
cp matrixMul ../bin/x86_64/linux/release
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 512
==20167== NVPROF is profiling process 20167, command: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 16
Grid Size X: 32, Grid Size Y: 32
Elapsed Time for CPU Multiplication: 1.480000 sec
SUCCESS!
==20167== Profiling application: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
==20167== Profiling result:
Time(%) Time Calls Avg Min Max Name
93.88% 30.588ms 1 30.588ms 30.588ms 30.588ms matrixMul(int*, int*, int*, int)
4.09% 1.3314ms 2 665.63us 662.70us 668.71us [CUDA memcpy HtoD]
2.03% 661.67us 1 661.67us 661.67us 661.67us [CUDA memcpy DtoH]

==20167== API calls:
Time(%) Time Calls Avg Min Max Name
96.29% 96.063ms 3 32.021ms 154.14us 95.742ms cudaMalloc
2.69% 2.6837ms 3 894.55us 290.04us 1.4712ms cudaMemcpy
0.58% 579.55us 83 6.9820us 813ns 235.62us cuDeviceGetAttribute
0.29% 285.91us 3 95.301us 83.644us 115.89us cudaFree
0.06% 55.484us 1 55.484us 55.484us 55.484us cuDeviceTotalMem
0.05% 47.024us 1 47.024us 47.024us 47.024us cuDeviceGetName
0.03% 33.415us 1 33.415us 33.415us 33.415us cudaLaunch
0.01% 9.8270us 4 2.4560us 889ns 6.7160us cudaSetupArgument
0.00% 4.1540us 2 2.0770us 872ns 3.2820us cuDeviceGet
0.00% 3.4100us 2 1.7050us 1.1490us 2.2610us cuDeviceGetCount
0.00% 2.7480us 1 2.7480us 2.7480us 2.7480us cudaConfigureCall
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$
```

Figure 8: Console Results for Matrix Multiplication Application Using Matrices of Size 512 x 512 and 16 Threads per Block

```
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul.o -c matrixMul.cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul matrixMul.o
mkdir -p ../bin/x86_64/linux/release
cp matrixMul ../bin/x86_64/linux/release
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 512
==20446== NVPROF is profiling process 20446, command: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 32
Grid Size X: 16, Grid Size Y: 16
Elapsed Time for CPU Multiplication: 1.490000 sec
SUCCESS!
==20446== Profiling application: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
==20446== Profiling result:
Time(%) Time Calls Avg Min Max Name
94.95% 36.696ms 1 36.696ms 36.696ms 36.696ms matrixMul(int*, int*, int*, int)
3.44% 1.3313ms 2 665.63us 662.70us 668.55us [CUDA memcpy HtoD]
1.61% 621.26us 1 621.26us 621.26us 621.26us [CUDA memcpy DtoH]

==20446== API calls:
Time(%) Time Calls Avg Min Max Name
96.03% 86.593ms 3 28.864ms 156.83us 86.270ms cudaMalloc
2.90% 2.6191ms 3 873.03us 289.99us 1.4063ms cudaMemcpy
0.58% 519.59us 83 6.2600us 904ns 209.95us cuDeviceGetAttribute
0.30% 275.00us 3 91.666us 78.147us 115.14us cudaFree
0.07% 59.411us 1 59.411us 59.411us 59.411us cuDeviceTotalMem
0.06% 49.817us 1 49.817us 49.817us 49.817us cuDeviceGetName
0.04% 34.343us 1 34.343us 34.343us 34.343us cudaLaunch
0.01% 10.223us 4 2.5550us 888ns 6.8470us cudaSetupArgument
0.00% 4.4690us 2 2.2340us 1.0110us 3.4580us cuDeviceGet
0.00% 3.7830us 2 1.8910us 1.3930us 2.3900us cuDeviceGetCount
0.00% 2.5620us 1 2.5620us 2.5620us 2.5620us cudaConfigureCall
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$
```

Figure 9: Console Results for Matrix Multiplication Application Using Matrices of Size 512 x 512 and 32 Threads per Block

```

[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul.o -c matrixMul.cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul matrixMul.o
mkdir -p ../bin/x86_64/linux/release
cp matrixMul ../bin/x86_64/linux/release
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 1024
==20754== NVPROF is profiling process 20754, command: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 2
Grid Size X: 512 , Grid Size Y: 512
Elapsed Time for CPU Multiplication: 10.860000 sec
SUCCESS!
==20754== Profiling application: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
==20754== Profiling result:
Time(%) Time Calls Avg Min Max Name
99.87% 5.97761s 1 5.97761s 5.97761s 5.97761s matrixMul(int*, int*, int*, int)
0.09% 5.3083ms 2 2.6542ms 2.6506ms 2.6578ms [CUDA memcpy HtoD]
0.04% 2.6866ms 1 2.6866ms 2.6866ms 2.6866ms [CUDA memcpy DtoH]

==20754== API calls:
Time(%) Time Calls Avg Min Max Name
90.63% 90.346ms 3 30.116ms 157.22us 90.022ms cudaMalloc
8.23% 8.2073ms 3 2.7358ms 1.8791ms 3.4044ms cudaMemcpy
0.57% 566.47us 83 6.8240us 1.1720us 208.60us cuDeviceGetAttribute
0.37% 364.75us 3 121.58us 87.632us 156.56us cudaFree
0.07% 70.449us 1 70.449us 70.449us 70.449us cuDeviceTotalMem
0.06% 57.249us 1 57.249us 57.249us 57.249us cuDeviceGetName
0.04% 44.603us 1 44.603us 44.603us 44.603us cudaLaunch
0.01% 14.007us 4 3.5010us 1.1250us 9.6940us cudaSetupArgument
0.01% 6.4580us 2 3.2290us 1.4660us 4.9920us cuDeviceGet
0.01% 5.3440us 2 2.6720us 2.0650us 3.2790us cuDeviceGetCount
0.00% 3.4230us 1 3.4230us 3.4230us 3.4230us cudaConfigureCall
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$

```

Figure 10: Console Results for Matrix Multiplication Application Using Matrices of Size 1024 x 1024 and 2 Threads per Block

```

[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul.o -c matrixMul.cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul matrixMul.o
mkdir -p ../bin/x86_64/linux/release
cp matrixMul ../bin/x86_64/linux/release
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 1024
==21029== NVPROF is profiling process 21029, command: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 4
Grid Size X: 256 , Grid Size Y: 256
Elapsed Time for CPU Multiplication: 10.930000 sec
SUCCESS!
==21029== Profiling application: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
==21029== Profiling result:
Time(%) Time Calls Avg Min Max Name
99.44% 1.54572s 1 1.54572s 1.54572s 1.54572s matrixMul(int*, int*, int*, int)
0.34% 5.3095ms 2 2.6547ms 2.6501ms 2.6594ms [CUDA memcpy HtoD]
0.22% 3.3651ms 1 3.3651ms 3.3651ms 3.3651ms [CUDA memcpy DtoH]

==21029== API calls:
Time(%) Time Calls Avg Min Max Name
91.38% 104.13ms 3 34.710ms 157.80us 103.80ms cudaMalloc
7.54% 8.5326ms 3 2.8642ms 1.8900ms 3.7996ms cudaMemcpy
0.54% 617.71us 83 7.4420us 1.1810us 223.98us cuDeviceGetAttribute
0.36% 414.45us 3 138.15us 105.40us 196.29us cudaFree
0.07% 77.172us 1 77.172us 77.172us 77.172us cuDeviceTotalMem
0.06% 63.035us 1 63.035us 63.035us 63.035us cuDeviceGetName
0.03% 34.641us 1 34.641us 34.641us 34.641us cudaLaunch
0.01% 10.804us 4 2.7010us 940ns 7.3930us cudaSetupArgument
0.01% 6.1470us 2 3.0730us 1.3920us 4.7550us cuDeviceGet
0.00% 5.0490us 2 2.5240us 1.8130us 3.2360us cuDeviceGetCount
0.00% 2.9970us 1 2.9970us 2.9970us 2.9970us cudaConfigureCall
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$

```

Figure 11: Console Results for Matrix Multiplication Application Using Matrices of Size 1024 x 1024 and 4 Threads per Block

```
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul.o -c matrixMul.cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul matrixMul.o
mkdir -p ../bin/x86_64/linux/release
cp matrixMul ../bin/x86_64/linux/release
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 1024
==21282== NVPROF is profiling process 21282, command: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 8
Grid Size X: 128 , Grid Size Y: 128
Elapsed Time for CPU Multiplication: 10.840000 sec
SUCCESS!
==21282== Profiling application: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
==21282== Profiling result:
Time(%) Time Calls Avg Min Max Name
98.08% 413.20ms 1 413.20ms 413.20ms 413.20ms matrixMul(int*, int*, int*, int)
1.26% 5.3088ms 2 2.6544ms 2.6502ms 2.6586ms [CUDA memcpy HtoD]
0.66% 2.7859ms 1 2.7859ms 2.7859ms 2.7859ms [CUDA memcpy DtoH]

==21282== API calls:
Time(%) Time Calls Avg Min Max Name
91.04% 96.304ms 3 32.101ms 156.80us 95.981ms cudaMalloc
7.88% 8.339ms 3 2.7800ms 1.8817ms 3.5358ms cudaMemcpy
0.56% 595.24us 83 7.1710us 1.1810us 224.88us cuDeviceGetAttribute
0.31% 329.06us 3 109.69us 82.784us 158.54us cudaFree
0.07% 77.070us 1 77.070us 77.070us 77.070us cuDeviceTotalMem
0.06% 62.349us 1 62.349us 62.349us 62.349us cuDeviceGetName
0.04% 47.332us 1 47.332us 47.332us 47.332us cudaLaunch
0.01% 13.327us 4 3.3310us 1.1670us 9.2150us cudaSetupArgument
0.01% 6.0580us 2 3.0290us 1.4350us 4.6230us cuDeviceGet
0.00% 4.9850us 2 2.4920us 1.9120us 3.0730us cuDeviceGetCount
0.00% 4.0040us 1 4.0040us 4.0040us 4.0040us cudaConfigureCall
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$
```

Figure 12: Console Results for Matrix Multiplication Application Using Matrices of Size 1024 x 1024 and 8 Threads per Block

```
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul.o -c matrixMul.cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul matrixMul.o
mkdir -p ../bin/x86_64/linux/release
cp matrixMul ../bin/x86_64/linux/release
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 1024
==21538== NVPROF is profiling process 21538, command: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 16
Grid Size X: 64 , Grid Size Y: 64
Elapsed Time for CPU Multiplication: 10.990000 sec
SUCCESS!
==21538== Profiling application: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
==21538== Profiling result:
Time(%) Time Calls Avg Min Max Name
96.87% 247.44ms 1 247.44ms 247.44ms 247.44ms matrixMul(int*, int*, int*, int)
2.08% 5.3076ms 2 2.6538ms 2.6503ms 2.6573ms [CUDA memcpy HtoD]
1.05% 2.6946ms 1 2.6946ms 2.6946ms 2.6946ms [CUDA memcpy DtoH]

==21538== API calls:
Time(%) Time Calls Avg Min Max Name
92.32% 112.55ms 3 37.517ms 153.72us 112.23ms cudaMalloc
6.66% 8.1196ms 3 2.7065ms 1.8810ms 3.3279ms cudaMemcpy
0.52% 636.93us 83 7.6730us 1.2600us 254.85us cuDeviceGetAttribute
0.32% 394.42us 3 131.47us 95.575us 198.85us cudaFree
0.07% 86.003us 1 86.003us 86.003us 86.003us cuDeviceTotalMem
0.05% 62.627us 1 62.627us 62.627us 62.627us cuDeviceGetName
0.03% 34.235us 1 34.235us 34.235us 34.235us cudaLaunch
0.01% 10.364us 4 2.5910us 947ns 6.9070us cudaSetupArgument
0.00% 6.0240us 2 3.0120us 1.3720us 4.6520us cuDeviceGet
0.00% 4.8630us 2 2.4340us 1.8360us 3.0330us cuDeviceGetCount
0.00% 2.8960us 1 2.8960us 2.8960us 2.8960us cudaConfigureCall
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$
```

Figure 13: Console Results for Matrix Multiplication Application Using Matrices of Size 1024 x 1024 and 16 Threads per Block

```

[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ make
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -I../common/inc -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul.o -c matrixMul.cu
/usr/local/cuda-7.0/bin/nvcc -ccbin g++ -m64 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -gencode arch=compute_37,code=sm_37 -gencode arch=compute_50,code=sm_50 -gencode arch=compute_52,code=sm_52 -gencode arch=compute_52,code=compute_52 -o matrixMul matrixMul.o
mkdir -p ../bin/x86_64/linux/release
cp matrixMul ../bin/x86_64/linux/release
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$ nvprof ~/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Square Matrix of Size: 1024
==21789== NVPROF is profiling process 21789, command: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
Threads per Block: 32
Grid Size X: 32 , Grid Size Y: 32
Elapsed Time for CPU Multiplication: 10.850000 sec
SUCCESS!
==21789== Profiling application: /home/student1/h2youssef/cuda-7.0/samples/0_Simple/matrixMul/matrixMul
==21789== Profiling result:
Time(%) Time Calls Avg Min Max Name
97.29% 288.65ms 1 288.65ms 288.65ms 288.65ms matrixMul(int*, int*, int*, int)
1.79% 5.3078ms 2 2.6539ms 2.6504ms 2.6574ms [CUDA memcpy HtoD]
0.92% 2.7228ms 1 2.7228ms 2.7228ms 2.7228ms [CUDA memcpy DtoH]

==21789== API calls:
Time(%) Time Calls Avg Min Max Name
92.62% 120.30ms 3 40.099ms 171.72us 119.95ms cudaMalloc
6.43% 8.3509ms 3 2.7836ms 1.9174ms 3.5110ms cudaMemcpy
0.50% 655.44us 83 7.8960us 1.1810us 283.74us cuDeviceGetAttribute
0.27% 357.11us 3 119.04us 94.592us 161.63us cudaFree
0.06% 77.321us 1 77.321us 77.321us 77.321us cuDeviceTotalMem
0.06% 76.484us 1 76.484us 76.484us 76.484us cuDeviceGetName
0.03% 37.633us 1 37.633us 37.633us 37.633us cudaLaunch
0.01% 11.653us 4 2.9130us 1.0620us 7.6060us cudaSetupArgument
0.00% 6.3890us 2 3.1940us 1.2540us 5.1350us cuDeviceGet
0.00% 4.9620us 2 2.4810us 1.8140us 3.1480us cuDeviceGetCount
0.00% 2.8010us 1 2.8010us 2.8010us 2.8010us cudaConfigureCall
[h2youssef@bloor:~/cuda-7.0/samples/0_Simple/matrixMul]$

```

Figure 13: Console Results for Matrix Multiplication Application Using Matrices of Size 1024 x 1024 and 16 Threads per Block